

EinsumsInC++

JUSTIN TURNEY

GitHub: github.com/jturney/EinsumsInC++

What is EinsumsInCpp?

Provides compile-time contraction pattern analysis to determine optimal tensor operation to perform.

```
std::vector<double> C(i * j);
std::vector<double> A(i * k);
std::vector<double> B(k * j);

for (size_t i = 0; i < i0; i++) {
    for (size_t j = 0; j < j0; j++) {
        for (size_t k = 0; k < k0; k++) {
            C[i * j0 + j] +=
                A[i * k0 + k] * B[k * j0 + j];
        }
    }
}
```



```
Tensor A{"A", i_, k_};
Tensor B{"B", k_, j_};
Tensor C{"C", i_, j_};

einsum(Indices{i, j}, &C,
        Indices{i, k}, A,
        Indices{k, j}, B);
```

Compile-Time Deduction

```
std::vector<double> C(i * j);  
std::vector<double> A(i * k);  
std::vector<double> B(k * j);  
  
for (size_t i = 0; i < i0; i++) {  
    for (size_t j = 0; j < j0; j++) {  
        for (size_t k = 0; k < k0; k++) {  
            C[i * j0 + j] +=  
                A[i * k0 + k] * B[k * j0 + j];  
        }  
    }  
}
```




```
Tensor A{"A", i_, k_};  
Tensor B{"B", k_, j_};  
Tensor C{"C", i_, j_};  
  
einsum(Indices{i, j}, &C,  
        Indices{i, k}, A,  
        Indices{k, j}, B);
```



GEMM


Compile-Time Deduction

```
Tensor A{"A", ....};  
Tensor B{"B", ....};  
Tensor C{"C", ....};  
  
einsum(Indices{....}, &C,  
       Indices{....}, A,  
       Indices{....}, B);
```



Compile-Time Deduction

```
Tensor A{"A", ....};  
Tensor B{"B", ....};  
Tensor C{"C", ....};  
  
einsum(CPrefactor, Indices{....}, &C,  
       ABPrefactor,  
       Indices{....}, A, Indices{....}, B);
```



How does it work?

```
Tensor A{"A", ....};  
Tensor B{"B", ....};  
Tensor C{"C", ....};
```

```
einsum(Indices{....}, &C,  
       Indices{....}, A,  
       Indices{....}, B);
```

```
template <typename AType, typename BType, typename CType, typename... CIndices, typename... AIndices,  
          typename... BIndices>  
auto einsum(const std::tuple<CIndices...> &C_indices, CType *C,  
            const std::tuple<AIndices...> &A_indices, const AType &A,  
            const std::tuple<BIndices...> &B_indices, const BType &B)  
-> std::enable_if_t<!is_smart_pointer_v<CType> && !is_smart_pointer_v<AType>  
                  && !is_smart_pointer_v<BType>> {  
    einsum(0, C_indices, C, 1, A_indices, A, B_indices, B);  
}
```

How does it work?

```
{
    einsum(0, C_indices, C, 1, A_indices, A, B_indices, B);
}

template <template <typename, size_t> typename AType, typename ADataType, size_t ARank,
          template <typename, size_t> typename BType, typename BDataType, size_t BRank,
          template <typename, size_t> typename CType, typename CDataType, size_t CRank,
          typename... CIndices, typename... AIndices, typename... BIndices>
auto einsum(const CDataType C_prefactor, const std::tuple<CIndices...> &C, CType<CDataType, CRank> *C,
            const std::conditional_t<(sizeof(ADataType) > sizeof(BDataType)), ADataType, BDataType> AB_prefactor,
            const std::tuple<AIndices...> &A, const AType<ADataType, ARank> &A,
            const std::tuple<BIndices...> &B, const BType<BDataType, BRank> &B)
-> std::enable_if_t<
    std::is_base_of_v<::einsums::detail::TensorBase<ADataType, ARank>, AType<ADataType, ARank>> &&
    std::is_base_of_v<::einsums::detail::TensorBase<BDataType, BRank>, BType<BDataType, BRank>> &&
    std::is_base_of_v<::einsums::detail::TensorBase<CDataType, CRank>, CType<CDataType, CRank>>>
{
```

How does it work?

It then performs a series of compile-time checks on the provided indices:

- Uniqueness
- Hadamard indices
- Index positions
- Contiguous indices
- And others.

```
constexpr auto A_indices = std::tuple<AIndices...>();
constexpr auto B_indices = std::tuple<BIndices...>();
constexpr auto C_indices = std::tuple<CIndices...>();
using ADataType = std::conditional_t<(sizeof(AIndices) > sizeof(BIndices)), AIndices, BIndices>;

// 1. Ensure the ranks are correct. (Compile-time check)
static_assert(sizeof...(CIndices) == CRank, "Rank of C does not match Indices given for C.");
static_assert(sizeof...(AIndices) == ARank, "Rank of A does not match Indices given for A.");
static_assert(sizeof...(BIndices) == BRank, "Rank of B does not match Indices given for B.");

// 2. Determine the links from AIndices and BIndices
constexpr auto linksAB = intersect_t<std::tuple<AIndices...>, std::tuple<BIndices...>>();
// 2a. Remove any links that appear in the target
constexpr auto links = difference_t<decltype(linksAB), std::tuple<CIndices...>>();

// 3. Determine the links between CIndices and AIndices
constexpr auto Clinks = intersect_t<std::tuple<CIndices...>, std::tuple<AIndices...>>();

// 4. Determine the links between CIndices and BIndices
constexpr auto Clinks = intersect_t<std::tuple<CIndices...>, std::tuple<BIndices...>>();

// Remove anything from A that exists in C
constexpr auto CminusA = difference_t<std::tuple<AIndices...>, std::tuple<AIndices...>>();
constexpr auto CminusB = difference_t<std::tuple<BIndices...>, std::tuple<BIndices...>>();

constexpr bool have_remaining_indices_in_CminusA = std::tuple_size_v<decltype(CminusA)>;
constexpr bool have_remaining_indices_in_CminusB = std::tuple_size_v<decltype(CminusB)>;

// Determine unique indices in A
constexpr auto A_only = difference_t<std::tuple<AIndices...>, decltype(links)>();
constexpr auto B_only = difference_t<std::tuple<BIndices...>, decltype(links)>();

constexpr auto A_unique = unique_t<std::tuple<AIndices...>>();
constexpr auto B_unique = unique_t<std::tuple<BIndices...>>();
constexpr auto C_unique = unique_t<std::tuple<CIndices...>>();
constexpr auto link_unique = c_unique_t<decltype(links)>();

constexpr bool A_hadamard_found = std::tuple_size_v<std::tuple<AIndices...>> != std::tuple_size_v<decltype(A_unique)>;
constexpr bool B_hadamard_found = std::tuple_size_v<std::tuple<BIndices...>> != std::tuple_size_v<decltype(B_unique)>;
constexpr bool C_hadamard_found = std::tuple_size_v<std::tuple<CIndices...>> != std::tuple_size_v<decltype(C_unique)>;

constexpr auto link_position_in_A = detail::find_type_with_position(link_unique, A_indices);
constexpr auto link_position_in_B = detail::find_type_with_position(link_unique, B_indices);
constexpr auto link_position_in_link = detail::find_type_with_position(link_unique, links);

constexpr auto target_position_in_A = detail::find_type_with_position(C_unique, A_indices);
constexpr auto target_position_in_B = detail::find_type_with_position(C_unique, B_indices);
constexpr auto target_position_in_C = detail::find_type_with_position(C_unique, C_indices);

constexpr auto A_target_position_in_C = detail::find_type_with_position(A_indices, C_indices);
constexpr auto B_target_position_in_C = detail::find_type_with_position(B_indices, C_indices);

auto unique_target_dims = detail::get_dim_ranges_for(*, detail::unique find_type_with_position(C_unique, C_indices));
auto unique_link_dims = detail::get_dim_ranges_for(A, link_position_in_A);

constexpr auto contiguous_link_position_in_A = detail::contiguous_positions(link_position_in_A);
constexpr auto contiguous_link_position_in_B = detail::contiguous_positions(link_position_in_B);

constexpr auto contiguous_target_position_in_A = detail::contiguous_positions(target_position_in_A);
constexpr auto contiguous_target_position_in_B = detail::contiguous_positions(target_position_in_B);

constexpr auto contiguous_A_targets_in_C = detail::contiguous_positions(A_target_position_in_C);
constexpr auto contiguous_B_targets_in_C = detail::contiguous_positions(B_target_position_in_C);

constexpr auto same_ordering_link_position_in_AB = detail::is_same_ordering(link_position_in_A, link_position_in_B);
constexpr auto same_ordering_target_position_in_CA = detail::is_same_ordering(target_position_in_A, A_target_position_in_C);
constexpr auto same_ordering_target_position_in_CB = detail::is_same_ordering(target_position_in_B, B_target_position_in_C);

constexpr auto C_exactly_matches_A =
    sizeof...(CIndices) == sizeof...(AIndices) && same_indices<std::tuple<CIndices...>, std::tuple<AIndices...>>();
constexpr auto C_exactly_matches_B =
    sizeof...(CIndices) == sizeof...(BIndices) && same_indices<std::tuple<CIndices...>, std::tuple<BIndices...>>();
constexpr auto A_exactly_matches_B = same_indices<std::tuple<AIndices...>, std::tuple<BIndices...>>();
```


How does it work?

```
// 2. Determine the links from AIndices and BIndices
constexpr auto linksAB = intersect_t<std::tuple<AIndices...>, std::tuple<BIndices...>>();
// 2a. Remove any links that appear in the target
constexpr auto links = difference_t<decltype(linksAB), std::tuple<CIndices...>>();

// 3. Determine the links between CIndices and AIndices
constexpr auto CALinks = intersect_t<std::tuple<CIndices...>, std::tuple<AIndices...>>();

// 4. Determine the links between CIndices and BIndices
constexpr auto CBLinks = intersect_t<std::tuple<CIndices...>, std::tuple<BIndices...>>();

// Remove anything from A that exists in C
constexpr auto CminusA = difference_t<std::tuple<CIndices...>, std::tuple<AIndices...>>();
constexpr auto CminusB = difference_t<std::tuple<CIndices...>, std::tuple<BIndices...>>();

constexpr bool have_remaining_indices_in_CminusA = std::tuple_size_v<decltype(CminusA)>;
constexpr bool have_remaining_indices_in_CminusB = std::tuple_size_v<decltype(CminusB)>;
```

How does it work?

```
constexpr auto is_gemm_possible = have_remaining_indices_in_CminusA && have_remaining_indices_in_CminusB &&
    contiguous_link_position_in_A && contiguous_link_position_in_B && contiguous_target_position_in_A &&
    contiguous_target_position_in_B && contiguous_A_targets_in_C && contiguous_B_targets_in_C &&
    same_ordering_link_position_in_AB && same_ordering_target_position_in_CA &&
    same_ordering_target_position_in_CB && !A_hadamard_found && !B_hadamard_found && !C_hadamard_found;
constexpr auto is_gemv_possible = contiguous_link_position_in_A && contiguous_link_position_in_B && contiguous_target_position_in_A &&
    same_ordering_link_position_in_AB && same_ordering_target_position_in_CA &&
    !same_ordering_target_position_in_CB && std::tuple_size_v<decltype(B_target_position_in_C)> == 0 &&
    !A_hadamard_found && !B_hadamard_found && !C_hadamard_found;

constexpr auto element_wise_multiplication =
    C_exactly_matches_A && C_exactly_matches_B && !A_hadamard_found && !B_hadamard_found && !C_hadamard_found;
constexpr auto dot_product =
    sizeof...(CIndices) == 0 && A_exactly_matches_B && !A_hadamard_found && !B_hadamard_found && !C_hadamard_found;

constexpr auto outer_product = std::tuple_size_v<decltype(linksAB)> == 0 && contiguous_target_position_in_A &&
    contiguous_target_position_in_B && !A_hadamard_found && !B_hadamard_found && !C_hadamard_found;
```

How does it work?

```
if constexpr (dot_product) {  
    CDataType temp = linear_algebra::dot(A, B);  
    (*C) *= C_prefactor;  
    (*C) += AB_prefactor * temp;  
  
    return;  
}
```

How does it work?

```
else if constexpr (element_wise_multiplication) {
    timer::Timer element_wise_multiplication{"element-wise multiplication"};

    auto target_dims = get_dim_ranges<CRank>(*C);
    auto view = std::apply(ranges::views::cartesian_product, target_dims);

    // Ensure the various tensors passed in are the same dimensionality
    if (((C->dims() != A.dims()) || C->dims() != B.dims())) {
        println_abort("einsum: at least one tensor does not have same dimensionality as destination");
    }

    #if defined(__INTEL_LLVM_COMPILER) || defined(__INTEL_COMPILER)
    #pragma omp parallel for simd
    #else
    #pragma omp parallel for
    #endif
    for (auto it = view.begin(); it != view.end(); it++) {
        CDataType &target_value = std::apply(*C, *it);
        ABDDataType AB_product = std::apply(A, *it) * std::apply(B, *it);
        target_value = C_prefactor * target_value + AB_prefactor * AB_product;
    }

    return;
}
```

How does it work?

If none of the patterns programmed can be utilized, then the code will use a generic, compile-time, threaded contraction code.

It's still better to work the equations or the tensors to match a better performing pattern.

```
einsum(1.0, Indices{p, q}, F, 1.0, Indices{p, q, r, s}, g, Indices{r, s}, D);  
einsum(1.0, Indices{p, q}, F, -1.0, Indices{p, r, q, s}, g, Indices{r, s}, D);
```

Sorting


```
sort(Indices{k, j, i}, &B,  
      Indices{i, j, k}, A);
```

```
sort(0.0, Indices{i, l, k, j}, &B,  
      0.5, Indices{k, j, l, i}, A);
```

Disk Tensors

EinsumsInC++ uses HDF5 library to store tensors to disk. It utilizes the h5cpp wrapper library to enable automatic tensor loading and storing to an HDF5 file.

HDF5 stores data into a hierarchal structure (think a file system within a file).



Disk Tensors

Storing the tensor properly on disk can be vitally important. EinsumsInC++ uses HDF5 which is sensitive to how the data is “chunked” on disk.

Timing information:

0 ms :	1000 calls :	0 ms per call	:	Timer Overhead
9206 ms :	1 calls :	9206 ms per call	:	Creating random tensor 64 64 64 64
0 ms :	1 calls :	0 ms per call	:	Creating disk tensor
2499 ms :	1 calls :	2499 ms per call	:	disk write (everything at once)
2401 ms :	1 calls :	2401 ms per call	:	disk read (equivalent to chunking)
2966 ms :	1 calls :	2966 ms per call	:	disk read (different to chunking)
0 ms :	1 calls :	0 ms per call	:	Creating disk tensor2
2605 ms :	1 calls :	2605 ms per call	:	disk write2 non-default chunk (everything at once)
2463 ms :	1 calls :	2463 ms per call	:	disk read2 (equivalent to chunking)
2938 ms :	1 calls :	2938 ms per call	:	disk read2 (different to chunking)
0 ms :	1 calls :	0 ms per call	:	Creating disk tensor3 bad chunking
17620 ms :	1 calls :	17620 ms per call	:	disk write3 (everything at once)
49799 ms :	1 calls :	49799 ms per call	:	disk read3 (different to chunking)

Disk Tensors

```
timer::push("Forming energy denominator (E_ijab)");
DiskTensor<double, 4> e_oovv{state::data,
    "/Method/Spin-Orbital/CCD/Energy Denominator oovv",
    nocc, nocc, nvir, nvir};

for (size_t i0 = 0; i0 < nocc; i0++) {
    for (size_t j0 = 0; j0 < nocc; j0++) {
        auto e_view = e_oovv(i0, j0, All, All);

        double e_ij = eocc(i0) + eocc(j0);
        for (size_t a0 = 0; a0 < nvir; a0++) {
            for (size_t b0 = 0; b0 < nvir; b0++) {
                e_view(a0, b0) =
                    1.0 / (e_ij - evir(a0) - evir(b0));
            }
        }
    }
}
timer::pop();
```

```
// Compute MP2 from the antisymmetrized integrals
timer::push("MP2 Check");
double e_mp2{0.0};
for (size_t i0 = 0; i0 < nocc; i0++) {
    for (size_t j0 = 0; j0 < nocc; j0++) {
        auto e_ij = e_oovv(i0, j0, All, All);
        auto g_ij = g_oovv(i0, j0, All, All);

        e_mp2 += linear_algebra::dot(g_ij.get(),
                                      g_ij.get(),
                                      e_ij.get());
    }
}
e_mp2 /= 4.0;
timer::pop();
```

Other Features

Perform contractions on any singular datatype and mixed datatypes.

Tensor Decompositions – Andy Jiang

- Parafac
- Tucker

Elemental Operations

- e.g.: abs, min, max, exp

Khatri-Rao Product

$$(A_{ij} \otimes B_{ij})_{ij}$$

Tensor Unfolding

$$A_{ijk} \rightarrow A_{i,\{jk\}}$$

$$A_{ijk} \rightarrow A_{j,\{ik\}}$$

Lyapunov Solver – Andy Jiang

$$AX + XA^T + Q = 0$$

Truncated SVD – Andy Jiang

Initial Implementation of Polynomials

- Gauss-Laguerre

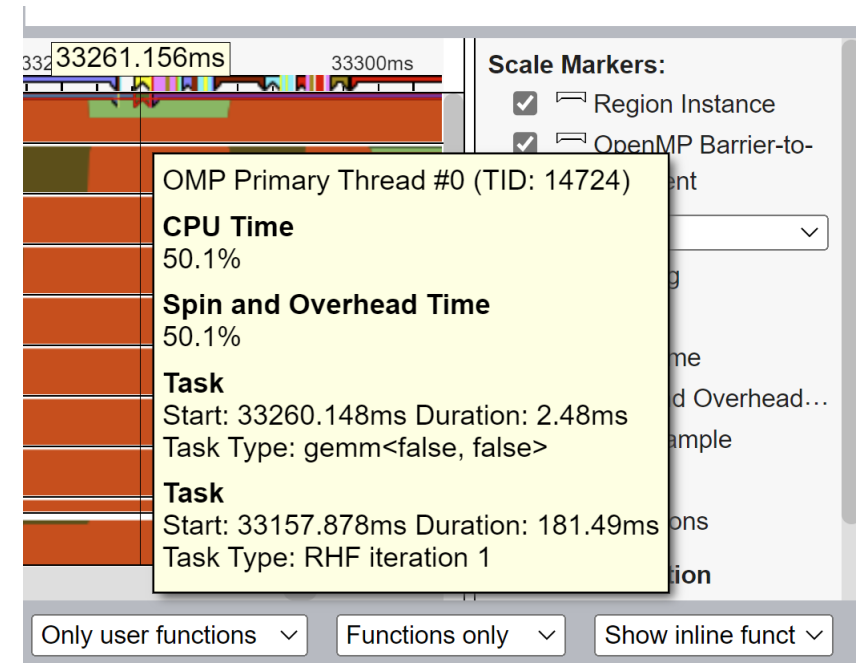
Plays Nice with VTune

If VTune is detected at CMake configure-time, the “Section” class will register with VTune to provide information on where in your code things are happening.

Instructions for running VTune in the EinsumsInC++ Docker container can be found [here](https://github.com/jturney/EinsumsInC++/tree/main/.devcontainer):

github.com/jturney/EinsumsInC++/tree/main/.devcontainer

```
{  
    Section section(fmt::format("gemm<{}, {}>", TransA, TransB));  
    blas::gemm(...);  
}
```



Where is it being used?

Tensor Hypercontraction Form of (T) Energy in Coupled-Cluster Theory – Andy Jiang

<https://arxiv.org/abs/2210.07035>

F12 Methods – Erica Mitchell

Future Plans

- Point group symmetry
- Identify new einsum patterns. E.g. batched gemms.
- Rudimentary auto-disk in einsum
- Interface to FFT
- Utilize oneMKL to provide GPU capabilities
- Incorporate TBLIS of D. Matthews as a backend for better CPU performance

GitHub: github.com/jturney/EinsumsInCpp